

REDUCE AND SCAN PATTERNS

A SHORT INTRODUCTION

Joseph Kehoe¹

¹Department of Computing and Networking
Institute of Technology Carlow

CDD101, 2017

TABLE OF CONTENTS

1 REDUCE

2 COMPLEX OPERATORS

3 SCAN

TABLE OF CONTENTS

1 REDUCE

2 COMPLEX OPERATORS

3 SCAN

DEFINITION

- The Reduce pattern is where a sequence of input data is reduced down to a single output value.
- A Combiner function is applied to every member of the input set.
- The combiner function operates a a pair of input values
result=combine(a,b)
- Examples include using $+$ to get the sum of a sequence
- We assume that the associated pairs of input can be combined in different orders and still get the same answers
- e.g. $a+b+c = b+a+c = b+c+a = a+c+b = c+a+b = c+b+a$

SIMPLE EXAMPLE

```
float sum(int dim, float in [])  
{  
    float sum=0.0;  
    for (int i=0; i < dim; ++i)  
    {  
        sum +=in[i];  
    }  
    return sum;  
}
```

SIMPLE OPENMP IMPLEMENTATION

Reduction is built in for simple operators

```
float sum(int dim, float in[])
{
    float sum=0.0;
    #pragma omp simd parallel for reduction(+:sum)
    for (int i=0; i < dim; ++i)
    {
        sum +=in[i];
    }
    return sum;
}
```

TABLE OF CONTENTS

1 REDUCE

2 COMPLEX OPERATORS

3 SCAN

MORE COMPLEX OPERATORS

- What if our operator is not supported by the **reduction** clause?
- Produce our own code using Tiling
- First each thread produces a result for its own subsequence
- Then combine all results for each tile into one value



SIMPLE OPENMP IMPLEMENTATION

```
float sum(int dim, float in[])
{
    float result=0.0;
    float tileResult[NumThreads];
    #pragma omp parallel for
    for (int i=0; i < dim; ++i)
    {
        int tid = omp_get_thread_num();
        tileResult[tid] =op(tileResult[tid], in[i]);
    }
    for (int i=0; i < NumThreads; ++i)
    {
        result=op(tileResult[i], result);
    }
    return result;
}
```

TABLE OF CONTENTS

1 REDUCE

2 COMPLEX OPERATORS

3 SCAN

DEFINITION

- Produces all partial reductions of an input sequence to produce a new output sequence
- used in e.g. integration
- Using sum as an example - each element will contain the sum of all previous elements
 - **Inclusive** scan means the sum includes everything up to and including the current element
 - **Exclusive** scan means we sum only the previous elements (do not include ourselves)
 - Show output of each type on sequence: 3, 4, 6, 8, 1, 4

APPROACH

If using standard fork join we use a three phase approach

- 1 Tile the sequence and reduce each tile in parallel
- 2 Do an exclusive scan of the reduction values (always exclusive)
- 3 Scan each of the tiles where the initial value is that calculated by the previous step

Or use **tasks** to implement a tree based approach