

# COMMUNICATION

## AN OVERVIEW

Joseph Kehoe<sup>1</sup>

<sup>1</sup>Department of Computing and Networking  
Institute of Technology Carlow

CDD101, 2018

# TABLE OF CONTENTS

- 1 OVERVIEW
- 2 PROTOCOLS
- 3 REMOTE PROCEDURE CALL
- 4 REMOTE OBJECT INVOCATION
- 5 MESSAGING

# TABLE OF CONTENTS

- 1 OVERVIEW
- 2 PROTOCOLS
- 3 REMOTE PROCEDURE CALL
- 4 REMOTE OBJECT INVOCATION
- 5 MESSAGING

# COMMUNICATION OVERVIEW

Computers in a distributed system may have different:

- Architectures
- Operating Systems
- Data encoding
- programming languages

We need a system of communication that allows them to talk to each other transparently

# TABLE OF CONTENTS

- 1 OVERVIEW
- 2 PROTOCOLS
- 3 REMOTE PROCEDURE CALL
- 4 REMOTE OBJECT INVOCATION
- 5 MESSAGING

Communication between computers is build on existing network protocols

- At the base we have TCP/IP or UDP/IP
- We can use socket programming to use these directly
- Or we can use higher level protocols instead of or alongside this: e.g. HTTP
- This base level protocol is available to all computers
- But we will need to add a layer on top of this base in order to give us the extra features we need

I am going to assume you are familiar with these lower level protocols

# TABLE OF CONTENTS

- 1 OVERVIEW
- 2 PROTOCOLS
- 3 REMOTE PROCEDURE CALL
- 4 REMOTE OBJECT INVOCATION
- 5 MESSAGING

Allow code to call procedures not resident on the local machine

- Causes a procedure to execute in a different address space (commonly on another computer on a shared network), which is coded as if it were a normal (local) procedure call, without the programmer explicitly coding the details for the remote interaction.
- The programmer writes essentially the same code whether the subroutine is local or remote
- Suited to Client Server Architectures
- RPC is a form of Inter Process Communication (IPC)
- RPC is usually synchronous in nature



## Sequence of events

- The client calls the client stub. The call is a local procedure call, with parameters pushed on to the stack in the normal way.
- The client stub packs the parameters into a message and makes a system call to send the message. Packing the parameters is called marshalling.
- The client's local operating system sends the message from the client machine to the server machine.
- The local operating system on the server machine passes the incoming packets to the server stub.
- The server stub unpacks the parameters from the message. Unpacking the parameters is called unmarshalling.
- Finally, the server stub calls the server procedure. The reply traces the same steps in the reverse direction.

To let different clients access servers, a number of standardized RPC systems have been created.

- Most of these use an interface description language (IDL) to let various platforms call the RPC
- The IDL files can then be used to generate code to interface between the client and servers.
- XML-RPC is an RPC protocol that uses XML to encode its calls and HTTP as a transport mechanism.
- JSON-RPC is an RPC protocol that uses JSON-encoded messages
- JSON-WSP is an RPC protocol that uses JSON-encoded messages
- SOAP is a successor of XML-RPC and also uses XML to encode its HTTP-based calls.

- Packaging of parameters for transport to remote server is called **Marshalling**
- Parameters are either:
  - Pass by Value
  - Pass by Reference
  - Pass by copy Restore
- Passing pointers is obviously very difficult

# TABLE OF CONTENTS

- 1 OVERVIEW
- 2 PROTOCOLS
- 3 REMOTE PROCEDURE CALL
- 4 REMOTE OBJECT INVOCATION
- 5 MESSAGING

- Extends RPC to objects
- Improves on RPC
- Objects separate state from interface (perfect for distributed Objects!)
- Keep state on server and put Interface on client
- Local proxy implements interface and holds reference to remote object
- It marshals data before sending, sends data, receives return values and unmarshalls them

# ROI (ON SERVER)

- Server has a server "stub" known as a skeleton
- This receives data, unmarshals it and calls the server object implementation. It then marshalls the return values and sends them back to the client proxy
- Often use the **Object Adapter** pattern to wrap an object interface around an existing system
- Object is transient if it cannot be offloaded from server. So if server goes down object is lost

# ROI (OBJECT REFERENCES)

- Object is persistent if it can be offloaded (saved) from the server to backup storage.
- This allows the object to be reanimated if the server crashes
- This even allows us to move the object around if we have a system wide object reference
- Need a DNS for object references if this is to work.

# ROI (DYNAMIC INVOCATION)

- If IDL is compiled at runtime then we have **Static Object Invocation**
- If we can decode interface at runtime then we have **Dynamic Object Invocation**
- Static Call is: *accountRef* –  $\rightarrow$  *deposit*(500);
- Dynamic Call is: *sendMessage(accountRef, deposit, 500)*



# ROI (EFFICIENCY)

- Remote Object Invocation can be inefficient
- Esp. if object state is small or parameters are large
- If system can distinguish between local and remote objects then we can circumvent this

- Integrated into the language (high level approach)
- Cloning of remote objects is difficult so local proxy is not cloned during process, just the remote object
- We must explicitly get a new reference to the remote cloned object
- Remote objects cannot be **synchronised** - only the proxy can be
- Why? hint: Client crashes during operation
- Anything serialisable can be marshalled
- Proxys can be serialised (so we can pass them around)

See here

# TABLE OF CONTENTS

- 1 OVERVIEW
- 2 PROTOCOLS
- 3 REMOTE PROCEDURE CALL
- 4 REMOTE OBJECT INVOCATION
- 5 MESSAGING

# SYNCHRONOUS AND ASYNCHRONOUS COMMUNICATION

**PERSISTENT ASYNCHRONOUS** Receiver need not be running when message is sent. Sender is not blocked

**PERSISTENT SYNCHRONOUS** Receiver need not be running when message is sent. Sender is blocked until acknowledgement is send from system receiving message (possibly before receiver starts running).

**TRANSIENT ASYNCHRONOUS** Sender is not blocke but can only send when receiver is running.

# SYNCHRONOUS AND ASYNCHRONOUS COMMUNICATION

**RECEIPT-BASED TRANSIENT SYNCHRONOUS** Receiver must be running for message to be sent and sender is blocked until receiver sends acknowledgement of receipt.

**DELIVERY-BASED TRANSIENT SYNCHRONOUS** Receiver must be running before message can be sent. Sender is blocked until receiver accepts message and starts working on it.

**RESPONSE-BASED TRANSIENT SYNCHRONOUS** Receiver must be running before message is sent. Sender is blocked until request is processed by receiver.

- Persistent Asynchronous Messaging
- Message can be send sucessfully even if receiving object/system is not be running at the time
- Sender sends and "forgets"
- Message transfer can take "minutes"
- Leads to loosely coupled systems
- Guarantees delivery but does not guarantee it will be "read"
- Suited to distributed systems (why?)
- SMTP is a good example (email)

## MESSAGE PASSING INTERFACE

- MPI is a standardized and portable message-passing standard designed by a group of researchers from academia and industry to function on a wide variety of parallel computing architectures
- The standard defines the syntax and semantics of a core of library routines in C, C++, Fortran and Python
- Commonly used on clusters



**MPI\_BSEND** Append outgoing message to local send buffer

**MPI\_SEND** Send a message and wait until copied to local or remote buffer

**MPI\_SSEND** Send a message and wait until receipt starts

**MPI\_SENDRECV** Send a message and wait for reply



# MPI PRIMITIVES

**MPI\_SEND** Pass reference to outgoing message, and continue

**MPI\_ISSEND** Pass refence to outgoing message , and wait until receipt starts

**MPI\_RECV** Receive a message, block if there is none

**MPI\_Irecv** Check if there is an incoming message, but do not block

see here